

Data Cleaning in R

John Brandt, Yale F&ES

March 15, 2018

Common Tasks

- ▶ Data class
- ▶ Creating dataframes
- ▶ Updating column names
- ▶ Combining columns
- ▶ Transposing
- ▶ Rownames/column names
- ▶ Concatenating
- ▶ Creating lists
- ▶ Subsetting (indices, logic, which)
- ▶ Merging
- ▶ Reclassification
- ▶ Dealing with NA values w/ logic
- ▶ Strings
- ▶ Functions
- ▶ lapply

Basic background

- ▶ `dataframe$column`
- ▶ `read.csv("data.csv")`
- ▶ `write.csv(object, "object.csv")`
- ▶ `head()`
- ▶ `str()`
- ▶ `colnames()`
- ▶ `summary()`
- ▶ `unique()`
- ▶ `class()`
- ▶ `levels()`
- ▶ `nrow()`, `length()`
- ▶ `dataframe[row, column]`

Structure

`as.Date`, `as.String`, `as.Numeric`, `as.Factor` use `class()` to find type

```
class(sales$event_dt)
```

```
[1] "factor"
```

```
sales$event_dt <- as.Date(sales$event_dt)
class(sales$event_dt)
```

```
[1] "Date"
```

Splitting columns

```
sales$event_dt[1:3]
```

```
[1] "2015-09-12" "2009-09-04" "2006-04-21"
```

```
strsplit(as.character(sales$event_dt[1:3]), "-")
```

```
[[1]]
```

```
[1] "2015" "09"   "12"
```

```
[[2]]
```

```
[1] "2009" "09"   "04"
```

```
[[3]]
```

```
[1] "2006" "04"   "21"
```

Creating new dataframes

```
new.df <- data.frame(matrix(nrow=5, ncol=3))
```

Updating column names

```
colnames(new.df) <- c("column.1", "column.2", "column.3")
new.df$column.1 <- c(1,2,3,4,5)
new.df$column.2 <- c("a", "b", "c", "d", "e")
new.df
```

	column.1	column.2	column.3
1	1	a	NA
2	2	b	NA
3	3	c	NA
4	4	d	NA
5	5	e	NA

Renaming singular column

```
colnames(new.df[colnames(new.df) == "column.2"])  
  <- "whatever"
```


Combining columns

```
new.df$column.3 <- paste(new.df$column.1,  
                          new.df$column.2, sep="")  
new.df$column.3
```

```
[1] "1a" "2b" "3c" "4d" "5e"
```

Transposing

`t()` returns a nested list. A dataframe must be specified if you want a dataframe.

```
new.df.t <- as.data.frame(t(new.df))  
new.df.t
```

	V1	V2	V3	V4	V5
column.1	1	2	3	4	5
column.2	a	b	c	d	e
column.3	1a	2b	3c	4d	5e

Rownames/colnames

```
require(tibble)  
rownames_to_column(new.df.t)
```

	rowname	V1	V2	V3	V4	V5
1	column.1	1	2	3	4	5
2	column.2	a	b	c	d	e
3	column.3	1a	2b	3c	4d	5e

Concatenating dataframes

```
new.df.2 <- new.df  
rbind(new.df, new.df.2)
```

	column.1	column.2	column.3
1	1	a	1a
2	2	b	2b
3	3	c	3c
4	4	d	4d
5	5	e	5e
6	1	a	1a
7	2	b	2b
8	3	c	3c
9	4	d	4d
10	5	e	5e

Creating new lists

```
mylist <- rep(NA, nrow(sales)/500)
```

```
mylist
```

```
[1] NA NA NA NA NA NA NA NA NA NA
```

Subsetting by indices

R makes use of the [row,column] notation

```
sales[5:10, c(3,9)]
```

	primary_act_id	major_cat_name
5	91c03a34b562436efa3c	MISC
6	ac4b847b3fde66f2117e	MISC
7	a14232befff04be1e2f3	MISC
8	91c03a34b562436efa3c	MISC
9	0efaba7ce3f0d7466b42	MISC
10	f6425a3223e73ea6de5a	CONCERTS

Subsetting by \geq , $==$, \leq , $!=$

R accepts logical statements within a [row, column] subsetting argument

```
sales[sales$age_yr>70 & !is.na(sales$age_yr),  
      c(34,14,15)]
```

	age_yr	tickets_purchased_qty	trans_face_val_amt
11	80	1	20
428	82	1	20
1220	80	1	20
1314	76	1	30
1701	72	2	30
3830	72	1	22
4188	90	1	15
4328	94	1	9
4975	78	1	20

Subsetting by %in%

```
sales[sales$venue_state %in% c("RHODE ISLAND",  
                               "MANITOBA"), c(18,26)]
```

	event_dt	venue_state
732	2015-08-08	RHODE ISLAND
1860	2015-09-17	MANITOBA
2046	2016-02-19	MANITOBA
2177	2015-12-16	RHODE ISLAND
2832	2015-09-17	MANITOBA
3304	2015-09-17	MANITOBA
4064	2016-01-27	MANITOBA
4119	2015-09-17	MANITOBA
4308	2015-11-16	MANITOBA
4935	2016-01-15	MANITOBA

Subsetting by which

If you have 500 columns, you may not know which column index to subset by in the previous example. Here the age column is extracted using which.

```
sales[sales$age_yr>70 & !is.na(sales$age_yr),  
      c(which(colnames(sales) %in%  
          c("age_yr","tickets_purchased_qty")))]
```

	tickets_purchased_qty	age_yr
11	1	80
428	1	82
1220	1	80
1314	1	76
1701	2	72
3830	1	72
4188	1	90
4328	1	94
4975	1	78

Which.max, which.min

`which.max` and `which.min` are useful for removing known outliers.

```
sales[-c(which.max(sales$income_amt)),]
```

Merging

```
merged <- merge(x1, x2, by.x="column.x", by.y="column.y")
```

Reclassifying

```
levels(cut(sales$trans_face_val_amt, 5))
```

```
[1] "(-0.52,305]"          "(305,609]"            "(609,913]"  
[4] "(913,1.22e+03]"       "(1.22e+03,1.52e+03]"
```

Removing NA values

```
new.df[2,3] <- NA  
new.df[1,2] <- NA  
new.df[1,3] <- NA  
new.df
```

	column.1	column.2	column.3
1	1	<NA>	<NA>
2	2	b	<NA>
3	3	c	3c
4	4	d	4d
5	5	e	5e

Na.omit

```
na.omit(new.df)
```

	column.1	column.2	column.3
3	3	c	3c
4	4	d	4d
5	5	e	5e

Column NAs

```
new.df[, colSums(is.na(new.df)) <= 1]
```

	column.1	column.2
1	1	<NA>
2	2	b
3	3	c
4	4	d
5	5	e

Row NAs

```
new.df[rowSums(is.na(new.df)) <= 1,]
```

	column.1	column.2	column.3
2	2	b	<NA>
3	3	c	3c
4	4	d	4d
5	5	e	5e

Converting date/time

lubridate package

```
year()  
month()  
day()  
week()  
as_date()  
as_datetime()  
time_length()
```

Lubridate examples

```
library(lubridate)
dates <- as.Date(sales$event_dt[1:3])
dates
```

```
[1] "2015-09-12" "2009-09-04" "2006-04-21"
```

```
year(dates)
```

```
[1] 2015 2009 2006
```

```
month(dates)
```

```
[1] 9 9 4
```

```
week(dates)
```

```
[1] 37 36 16
```

```
day(dates)
```

More examples

```
floor_date(dates, "month")
```

```
[1] "2015-09-01" "2009-09-01" "2006-04-01"
```

```
ceiling_date(dates, "season")
```

```
[1] "2015-12-01" "2009-12-01" "2006-06-01"
```

```
date1 <- "2009-08-03 12:01:59"  
as.Date(date1) # uh-oh
```

```
[1] "2009-08-03"
```

```
as_datetime(date1)
```

```
[1] "2009-08-03 12:01:59 UTC"
```

Strings

grepl returns a logical TRUE/FALSE

```
files <- files[grepl('name', files) == TRUE]
```

```
gsub("hello","goodbye",files)
```

Creating functions

```
myfunction <- function(x) {  
  z <- x + 1  
  return(z)  
}
```

```
myfunction(3)
```

```
[1] 4
```

Function to calculate percentage of NA in columns

```
calc.na <- function(x, data) {  
  calc <- sum(is.na(data[[x]]))/nrow(data)  
  return(unlist(calc))  
}  
  
calc.na("onsale_dt", sales)
```

```
[1] 0.0202
```

What about the amount of every column?

Option 1: For loop

Problem 1: results duplicated Problem 2: returns list indices as column names, which cannot be iterated over or used in the future.

```
results <- rep(NA, length(colnames(sales)))
```

```
for (i in colnames(sales)) {  
  results[i] <- calc.na(i, sales)  
}
```

```
results[60]
```

```
tickets_purchased_qty  
0
```

Option 2: lapply

```
unlist(lapply(colnames(sales), calc.na, sales))
```

```
[1] 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.682  
[11] 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.000  
[21] 0.0008 0.0000 0.0848 0.0000 0.0000 0.0000 0.0000 0.484  
[31] 0.0000 0.0000 0.9646 0.9664 0.9692 0.9722 0.9776 0.982  
[41] 0.9740 0.9724 0.9842 0.9846 0.9874 0.9354
```


Function - example 2

```
toMatch <- c("x", "y", "z", "...")
sentences <- c("sentence 1", "sentence 2", "...")

subset_sentences <- function(Match, sentences){
  sentences[grep(Match,sentences)]
}

subsetting <- lapply(toMatch, subset_sentences, sentences)
```

Read in data

```
require(dplyr)
require(lubridate)
sales <- read.csv("Workshops/sales.csv",
                  na.strings=c("NULL", "NA"))
```

Function

```
perc.na <- function(data) {  
  sing.na <- function(x) {  
    sum(is.na(data[[x]])/nrow(data))  
  }  
  return(unlist(lapply(colnames(data), sing.na)))  
}  
  
sales.na <- perc.na(sales)
```

Remove outliers, select columns

```
sales <- sales[, -which(sales.na >= 0.9)]
```

```
sales <- sales %>%  
  select(-c(X, primary_act_id, secondary_act_id,  
            primary_act_id, secondary_act_id,  
            purch_party_lkup_id,  
            web_session_cookie_val))
```

Reclassify

```
sales$secondary <- 0  
sales$secondary[!is.na(sales$secondary_act_name)] <- 1  
sales$secondary <- as.factor(sales$secondary)
```

Calc new columns

```
sales$onsale_dt <- as.Date(sales$onsale_dt)
sales$event_dt <- as.Date(sales$event_dt)
sales$sales_ord_tran_dt <-
  as.Date(sales$sales_ord_tran_dt)

sales$sale_length <- sales$event_dt - sales$onsale_dt
sales$purchased_date <- sales$sales_ord_tran_dt -
  sales$onsale_dt

sales$event_hour <- hour(sales$event_date_time)
```

Remove unnecessary columns before aggregating

```
sales <- sales[,-c(1,2,3,4,7,8,9,13,  
                  15,16,17,18,19,20,  
                  23,24,25,26,27)]
```

Datetime conversions

```
sales$month <- month(sales$event_dt)
sales$year <- year(sales$event_dt)
sales$month_year <- paste(sales$month, sales$year)
```


dplyr for final data

Go to Matt's data exploration workshop to learn more about this!

```
sales.y <- sales %>%  
  select(-c(minor_cat_name, venue_city, event_dt,  
            month, month_year, secondary,  
            major_cat_name)) %>%  
  group_by(year, delivery_type_cd,  
            venue_state) %>%  
  summarise_all(c("mean", "sum"))
```